

# RTML 2025 | A1 | Pytorch-AlexNet-GoogleNet

## Final Report

---

Name: Arya Shah

StudentID: st125462

This notebook consists of solutions for the following exercises:

(⚠ Note that this is a Jupyter Notebook compiled from the individual python scripts and is not intended to be run as is)

1. Create these three networks. Be sure to properly define your Python classes, with one class per file and a main module that sets up your objects, runs the training process, and saves the necessary data.
2. Note that the AlexNet implementation here does not have the local response normalization feature described in the paper. Take a look at the PyTorch implementation of LRN and incorporate it into your AlexNet implementation as it is described in the paper. Compare your test set results with and without LRN.
3. Note that the backbone of the GoogLeNet implemented thus far does not correspond exactly to the description. Modify the architecture to Use the same backbone (input image size, convolutions, etc.) before the first Inception module  
Add the two side classifiers
4. Compare your GoogLeNet and AlexNet implementations on CIFAR-10. Comment on the number of parameters, speed of training, and accuracy of the two models on this dataset when trained from scratch.
5. Experiment with the pretrained GoogLeNet and AlexNet from the torchvision repository. Does it give better results on CIFAR-10 similar to what we found with AlexNet? Comment on what we can glean from the results about the capacity and generalization ability of these two models.

## AlexNet Class

✅ SOLVED: Note that the AlexNet implementation here does not have the local response normalization feature described in the paper. Take a look at the PyTorch implementation of LRN and incorporate it into your AlexNet implementation as it is described in the paper. Compare your test set results with and without LRN.

```
In [ ]: # For our puffer server we need to browse via a proxy!!  
import os
```

```

# Set HTTP and HTTPS proxy
os.environ['http_proxy'] = 'http://192.41.170.23:3128'
os.environ['https_proxy'] = 'http://192.41.170.23:3128'

import torch
import torch.nn as nn

class AlexNet(nn.Module):
    """
    An AlexNet-like CNN with Local Response Normalization (LRN)

    Attributes
    -----
    num_classes : int
        Number of classes in the final multinomial output layer
    features : Sequential
        The feature extraction portion of the network
    avgpool : AdaptiveAvgPool2d
        Convert the final feature layer to 6x6 feature maps by average pooling
    classifier : Sequential
        Classify the feature maps into num_classes classes
    use_lrn : bool
        Whether to use Local Response Normalization
    """
    def __init__(self, num_classes: int = 10, use_lrn: bool = True):
        super().__init__()
        self.num_classes = num_classes
        self.use_lrn = use_lrn

        # First conv layer
        self.features_1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
            nn.ReLU(inplace=True),
        )

        # First LRN layer (after ReLU, before MaxPool)
        self.lrn1 = nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2)

        # First MaxPool and second conv
        self.features_2 = nn.Sequential(
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(64, 192, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
        )

        # Second LRN layer (after ReLU, before MaxPool)
        self.lrn2 = nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2)

        # Second MaxPool and remaining layers
        self.features_3 = nn.Sequential(
            nn.MaxPool2d(kernel_size=3, stride=2),
            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )

```

```

self.avgpool = nn.AdaptiveAvgPool2d((6, 6))
self.classifier = nn.Sequential(
    nn.Dropout(p=0.5), # Added dropout rate as per paper
    nn.Linear(256 * 6 * 6, 4096),
    nn.ReLU(inplace=True),
    nn.Dropout(p=0.5), # Added dropout rate as per paper
    nn.Linear(4096, 4096),
    nn.ReLU(inplace=True),
    nn.Linear(4096, num_classes),
)

def forward(self, x: torch.Tensor) -> torch.Tensor:
    # First conv + ReLU
    x = self.features_1(x)

    # First LRN (if enabled)
    if self.use_lrn:
        x = self.lrn1(x)

    # First MaxPool + second conv + ReLU
    x = self.features_2(x)

    # Second LRN (if enabled)
    if self.use_lrn:
        x = self.lrn2(x)

    # Remaining Layers
    x = self.features_3(x)
    x = self.avgpool(x)
    x = torch.flatten(x, 1)
    x = self.classifier(x)
    return x

```

## Output & Results:

### 1. Training AlexNet With LRN

```

Training AlexNet with LRN...
Epoch 1/20:
Train Loss: 2.3026, Train Acc: 0.1009
Val Loss: 2.3020, Val Acc: 0.1176
Batch Time: 0.0353s
Epoch 2/20:
Train Loss: 2.1474, Train Acc: 0.2016
Val Loss: 1.9162, Val Acc: 0.2708
Batch Time: 0.0381s
Epoch 3/20:
Train Loss: 1.7291, Train Acc: 0.3462
Val Loss: 1.5222, Val Acc: 0.4324
Batch Time: 0.0357s
Epoch 4/20:
Train Loss: 1.4334, Train Acc: 0.4713
Val Loss: 1.3261, Val Acc: 0.5064
Batch Time: 0.0357s
Epoch 5/20:
Train Loss: 1.2513, Train Acc: 0.5426

```

Val Loss: 1.1364, Val Acc: 0.5862  
Batch Time: 0.0361s  
Epoch 6/20:  
Train Loss: 1.0832, Train Acc: 0.6118  
Val Loss: 0.9491, Val Acc: 0.6574  
Batch Time: 0.0367s  
Epoch 7/20:  
Train Loss: 0.9441, Train Acc: 0.6650  
Val Loss: 0.8691, Val Acc: 0.6924  
Batch Time: 0.0356s  
Epoch 8/20:  
Train Loss: 0.8381, Train Acc: 0.7050  
Val Loss: 0.8000, Val Acc: 0.7080  
Batch Time: 0.0358s  
Epoch 9/20:  
Train Loss: 0.7660, Train Acc: 0.7301  
Val Loss: 0.6960, Val Acc: 0.7560  
Batch Time: 0.0348s  
Epoch 10/20:  
Train Loss: 0.6950, Train Acc: 0.7567  
Val Loss: 0.6654, Val Acc: 0.7678  
Batch Time: 0.0361s  
Epoch 11/20:  
Train Loss: 0.6475, Train Acc: 0.7753  
Val Loss: 0.6335, Val Acc: 0.7772  
Batch Time: 0.0403s  
Epoch 12/20:  
Train Loss: 0.5952, Train Acc: 0.7919  
Val Loss: 0.6191, Val Acc: 0.7792  
Batch Time: 0.0385s  
Epoch 13/20:  
Train Loss: 0.5538, Train Acc: 0.8068  
Val Loss: 0.5497, Val Acc: 0.8060  
Batch Time: 0.0389s  
Epoch 14/20:  
Train Loss: 0.5155, Train Acc: 0.8229  
Val Loss: 0.5501, Val Acc: 0.8102  
Batch Time: 0.0377s  
Epoch 15/20:  
Train Loss: 0.4816, Train Acc: 0.8315  
Val Loss: 0.5513, Val Acc: 0.8054  
Batch Time: 0.0376s  
Epoch 16/20:  
Train Loss: 0.4449, Train Acc: 0.8455  
Val Loss: 0.5023, Val Acc: 0.8252  
Batch Time: 0.0407s  
Epoch 17/20:  
Train Loss: 0.4196, Train Acc: 0.8529  
Val Loss: 0.5053, Val Acc: 0.8260  
Batch Time: 0.0371s  
Epoch 18/20:  
Train Loss: 0.3952, Train Acc: 0.8619  
Val Loss: 0.5119, Val Acc: 0.8234  
Batch Time: 0.0426s  
Epoch 19/20:  
Train Loss: 0.3698, Train Acc: 0.8722

Val Loss: 0.4756, Val Acc: 0.8362  
Batch Time: 0.0422s  
Epoch 20/20:  
Train Loss: 0.3515, Train Acc: 0.8784  
Val Loss: 0.4513, Val Acc: 0.8446  
Batch Time: 0.0367s  
Total training time for AlexNet with LRN: 17.12 minutes

## 2. Training Alexnet Without LRN

Training AlexNet without LRN...

Epoch 1/20:  
Train Loss: 2.2073, Train Acc: 0.1632  
Val Loss: 1.9603, Val Acc: 0.2782  
Batch Time: 0.0383s

Epoch 2/20:  
Train Loss: 1.6577, Train Acc: 0.3869  
Val Loss: 1.4239, Val Acc: 0.4726  
Batch Time: 0.0405s

Epoch 3/20:  
Train Loss: 1.3588, Train Acc: 0.5036  
Val Loss: 1.2013, Val Acc: 0.5752  
Batch Time: 0.0415s

Epoch 4/20:  
Train Loss: 1.1139, Train Acc: 0.6011  
Val Loss: 1.0265, Val Acc: 0.6304  
Batch Time: 0.0414s

Epoch 5/20:  
Train Loss: 0.9415, Train Acc: 0.6672  
Val Loss: 0.8589, Val Acc: 0.6946  
Batch Time: 0.0432s

Epoch 6/20:  
Train Loss: 0.8340, Train Acc: 0.7078  
Val Loss: 0.8179, Val Acc: 0.7124  
Batch Time: 0.0424s

Epoch 7/20:  
Train Loss: 0.7398, Train Acc: 0.7423  
Val Loss: 0.7249, Val Acc: 0.7478  
Batch Time: 0.0430s

Epoch 8/20:  
Train Loss: 0.6628, Train Acc: 0.7690  
Val Loss: 0.6563, Val Acc: 0.7736  
Batch Time: 0.0385s

Epoch 9/20:  
Train Loss: 0.6059, Train Acc: 0.7881  
Val Loss: 0.5994, Val Acc: 0.7886  
Batch Time: 0.0405s

Epoch 10/20:  
Train Loss: 0.5562, Train Acc: 0.8072  
Val Loss: 0.6072, Val Acc: 0.7878  
Batch Time: 0.0395s

Epoch 11/20:  
Train Loss: 0.5121, Train Acc: 0.8240  
Val Loss: 0.5246, Val Acc: 0.8150  
Batch Time: 0.0433s

Epoch 12/20:  
Train Loss: 0.4741, Train Acc: 0.8329  
Val Loss: 0.5355, Val Acc: 0.8130  
Batch Time: 0.0428s  
Epoch 13/20:  
Train Loss: 0.4421, Train Acc: 0.8485  
Val Loss: 0.5294, Val Acc: 0.8144  
Batch Time: 0.0426s  
Epoch 14/20:  
Train Loss: 0.4087, Train Acc: 0.8582  
Val Loss: 0.5418, Val Acc: 0.8152  
Batch Time: 0.0427s  
Epoch 15/20:  
Train Loss: 0.3846, Train Acc: 0.8661  
Val Loss: 0.4777, Val Acc: 0.8366  
Batch Time: 0.0405s  
Epoch 16/20:  
Train Loss: 0.3459, Train Acc: 0.8802  
Val Loss: 0.4799, Val Acc: 0.8408  
Batch Time: 0.0414s  
Epoch 17/20:  
Train Loss: 0.3303, Train Acc: 0.8858  
Val Loss: 0.4991, Val Acc: 0.8304  
Batch Time: 0.0362s  
Epoch 18/20:  
Train Loss: 0.3070, Train Acc: 0.8925  
Val Loss: 0.4683, Val Acc: 0.8492  
Batch Time: 0.0408s  
Epoch 19/20:  
Train Loss: 0.2867, Train Acc: 0.9004  
Val Loss: 0.4610, Val Acc: 0.8506  
Batch Time: 0.0365s  
Epoch 20/20:  
Train Loss: 0.2718, Train Acc: 0.9049  
Val Loss: 0.4816, Val Acc: 0.8374  
Batch Time: 0.0437s  
Total training time for AlexNet without LRN: 17.47 minutes

## Key Observations:

### I. LRN Implementation Details

- Two LRN layers are implemented with parameters:
- size: 5 (local neighborhood)
- alpha: 0.0001 (scaling parameter)
- beta: 0.75 (exponent)
- k: 2 (additive constant)

### II. LRN Placement

- First LRN layer: After the first ReLU and before MaxPool
- Second LRN layer: After the second ReLU and before MaxPool

- The implementation uses a modular approach with three feature blocks. LRN layers are conditionally applied using the `use_lrn` flag
1. The version with and without LRN did not demonstrate any major improvements:
    - Final test accuracy quite similar (84.56% vs 83.78%)
    - Better training accuracy for without LRN (90.49% vs 87.84%)
  2. Training stability:
    - Both models showed similar batch processing times
    - Training time difference was minimal (0.35 minutes)
  3. Generalization:
    - Both models showed good generalization with minimal overfitting

## GoogleNet Class

✓ SOLVED: Note that the backbone of the GoogLeNet implemented thus far does not correspond exactly to the description. Modify the architecture to Use the same backbone (input image size, convolutions, etc.) before the first Inception module Add the two side classifiers

```
In [ ]: # For our puffer server we need to browse via a proxy!!
import os
# Set HTTP and HTTPS proxy
os.environ['http_proxy'] = 'http://192.41.170.23:3128'
os.environ['https_proxy'] = 'http://192.41.170.23:3128'

import torch
import torch.nn as nn
import torch.nn.functional as F

class Inception(nn.Module):
    """
    Inception block for a GoogLeNet-like CNN

    Attributes
    -----
    in_planes : int
        Number of input feature maps
    n1x1 : int
        Number of direct 1x1 convolutions
    n3x3red : int
        Number of 1x1 reductions before the 3x3 convolutions
    n3x3 : int
        Number of 3x3 convolutions
    n5x5red : int
        Number of 1x1 reductions before the 5x5 convolutions
    n5x5 : int
        Number of 5x5 convolutions
    pool_planes : int
        Number of 1x1 convolutions after 3x3 max pooling
```

```

...
def __init__(self, in_planes, n1x1, n3x3red, n3x3, n5x5red, n5x5, pool_planes):
    super(Inception, self).__init__()
    self.in_planes = in_planes
    self.n1x1 = n1x1
    self.n3x3red = n3x3red
    self.n3x3 = n3x3
    self.n5x5red = n5x5red
    self.n5x5 = n5x5
    self.pool_planes = pool_planes

    # 1x1 conv branch
    self.b1 = nn.Sequential(
        nn.Conv2d(in_planes, n1x1, kernel_size=1),
        nn.BatchNorm2d(n1x1),
        nn.ReLU(True),
    )

    # 1x1 conv -> 3x3 conv branch
    self.b2 = nn.Sequential(
        nn.Conv2d(in_planes, n3x3red, kernel_size=1),
        nn.BatchNorm2d(n3x3red),
        nn.ReLU(True),
        nn.Conv2d(n3x3red, n3x3, kernel_size=3, padding=1),
        nn.BatchNorm2d(n3x3),
        nn.ReLU(True),
    )

    # 1x1 conv -> 5x5 conv branch
    self.b3 = nn.Sequential(
        nn.Conv2d(in_planes, n5x5red, kernel_size=1),
        nn.BatchNorm2d(n5x5red),
        nn.ReLU(True),
        nn.Conv2d(n5x5red, n5x5, kernel_size=3, padding=1),
        nn.BatchNorm2d(n5x5),
        nn.ReLU(True),
        nn.Conv2d(n5x5, n5x5, kernel_size=3, padding=1),
        nn.BatchNorm2d(n5x5),
        nn.ReLU(True),
    )

    # 3x3 pool -> 1x1 conv branch
    self.b4 = nn.Sequential(
        nn.MaxPool2d(3, stride=1, padding=1),
        nn.Conv2d(in_planes, pool_planes, kernel_size=1),
        nn.BatchNorm2d(pool_planes),
        nn.ReLU(True),
    )

    def forward(self, x):
        y1 = self.b1(x)
        y2 = self.b2(x)
        y3 = self.b3(x)
        y4 = self.b4(x)
        return torch.cat([y1, y2, y3, y4], 1)

class InceptionAux(nn.Module):
    ...
    Auxiliary classifier for GoogLeNet

```



```

Attributes
-----
conv : Sequential
    Convolutional layers for feature extraction
fc1 : Linear
    First fully connected layer
fc2 : Linear
    Output layer
dropout : Dropout
    Dropout layer for regularization
...

def __init__(self, in_channels, num_classes):
    super(InceptionAux, self).__init__()
    self.conv = nn.Sequential(
        nn.AvgPool2d(kernel_size=5, stride=3),
        nn.Conv2d(in_channels, 128, kernel_size=1),
        nn.ReLU(True)
    )
    self.fc1 = nn.Linear(128 * 4 * 4, 1024)
    self.fc2 = nn.Linear(1024, num_classes)
    self.dropout = nn.Dropout(0.7)

def forward(self, x):
    x = self.conv(x)
    x = torch.flatten(x, 1)
    x = F.relu(self.fc1(x))
    x = self.dropout(x)
    x = self.fc2(x)
    return x

class GoogLeNet(nn.Module):
    ...

    GoogLeNet (Inception v1) implementation

Attributes
-----
pre_layers : Sequential
    Initial convolutional layers before inception modules
a3-b5 : Inception
    Inception blocks
aux1, aux2 : InceptionAux
    Auxiliary classifiers
avgpool : AvgPool2d
    Average pool layer after final inception block
dropout : Dropout
    Dropout layer before final classifier
fc : Linear
    Final classifier layer
...

def __init__(self, num_classes=10, aux_logits=True, transform_input=False):
    super(GoogLeNet, self).__init__()
    self.aux_logits = aux_logits
    self.transform_input = transform_input

    # Initial layers before inception modules (matching paper)
    self.conv1 = nn.Sequential(
        nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3),
        nn.ReLU(True),
        nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
        nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2)
    )

```

```

)

self.conv2 = nn.Sequential(
    nn.Conv2d(64, 64, kernel_size=1),
    nn.ReLU(True),
    nn.Conv2d(64, 192, kernel_size=3, padding=1),
    nn.ReLU(True),
    nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2),
    nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
)

self.inception3a = Inception(192, 64, 96, 128, 16, 32, 32)
self.inception3b = Inception(256, 128, 128, 192, 32, 96, 64)
self.maxpool3 = nn.MaxPool2d(3, stride=2, padding=1)

self.inception4a = Inception(480, 192, 96, 208, 16, 48, 64)
self.inception4b = Inception(512, 160, 112, 224, 24, 64, 64)
self.inception4c = Inception(512, 128, 128, 256, 24, 64, 64)
self.inception4d = Inception(512, 112, 144, 288, 32, 64, 64)
self.inception4e = Inception(528, 256, 160, 320, 32, 128, 128)
self.maxpool4 = nn.MaxPool2d(3, stride=2, padding=1)

self.inception5a = Inception(832, 256, 160, 320, 32, 128, 128)
self.inception5b = Inception(832, 384, 192, 384, 48, 128, 128)

if aux_logits:
    self.aux1 = InceptionAux(512, num_classes)
    self.aux2 = InceptionAux(528, num_classes)

self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
self.dropout = nn.Dropout(0.4)
self.fc = nn.Linear(1024, num_classes)

def _transform_input(self, x):
    if self.transform_input:
        x_ch0 = torch.unsqueeze(x[:, 0], 1) * (0.229 / 0.5) + (0.485 - 0.5)
        x_ch1 = torch.unsqueeze(x[:, 1], 1) * (0.224 / 0.5) + (0.456 - 0.5)
        x_ch2 = torch.unsqueeze(x[:, 2], 1) * (0.225 / 0.5) + (0.406 - 0.5)
        x = torch.cat((x_ch0, x_ch1, x_ch2), 1)
    return x

def forward(self, x):
    # N x 3 x 224 x 224
    x = self._transform_input(x)

    # N x 64 x 112 x 112
    x = self.conv1(x)

    # N x 192 x 56 x 56
    x = self.conv2(x)

    # N x 256 x 56 x 56
    x = self.inception3a(x)
    # N x 480 x 56 x 56
    x = self.inception3b(x)
    # N x 480 x 28 x 28
    x = self.maxpool3(x)

    # N x 512 x 28 x 28
    x = self.inception4a(x)

```

```

        if self.training and self.aux_logits:
            aux1 = self.aux1(x)

        x = self.inception4b(x)
        x = self.inception4c(x)
        x = self.inception4d(x)

        if self.training and self.aux_logits:
            aux2 = self.aux2(x)

        x = self.inception4e(x)
        x = self.maxpool4(x)

        x = self.inception5a(x)
        x = self.inception5b(x)

        x = self.avgpool(x)
        x = torch.flatten(x, 1)
        x = self.dropout(x)
        x = self.fc(x)

        if self.training and self.aux_logits:
            return x, aux1, aux2
        return x

```

## Output & Results

```

Training Custom GoogLeNet...
Epoch 1/20:
Train Loss: 2.4877, Train Acc: 0.4582
Val Loss: 1.3349, Val Acc: 0.5264
Batch Time: 0.0830s
Epoch 2/20:
Train Loss: 1.6215, Train Acc: 0.6679
Val Loss: 1.0407, Val Acc: 0.6516
Batch Time: 0.0821s
Epoch 3/20:
Train Loss: 1.2190, Train Acc: 0.7613
Val Loss: 0.7700, Val Acc: 0.7368
Batch Time: 0.0824s
Epoch 4/20:
Train Loss: 0.9855, Train Acc: 0.8140
Val Loss: 0.6541, Val Acc: 0.7804
Batch Time: 0.0793s
Epoch 5/20:
Train Loss: 0.8510, Train Acc: 0.8427
Val Loss: 0.5375, Val Acc: 0.8190
Batch Time: 0.0818s
Epoch 6/20:
Train Loss: 0.7347, Train Acc: 0.8696
Val Loss: 0.5920, Val Acc: 0.8028
Batch Time: 0.0865s
Epoch 7/20:
Train Loss: 0.6557, Train Acc: 0.8856
Val Loss: 0.6266, Val Acc: 0.7988

```

Batch Time: 0.0869s  
Epoch 8/20:  
Train Loss: 0.5879, Train Acc: 0.8993  
Val Loss: 0.5096, Val Acc: 0.8342  
Batch Time: 0.0787s  
Epoch 9/20:  
Train Loss: 0.5376, Train Acc: 0.9102  
Val Loss: 0.4639, Val Acc: 0.8470  
Batch Time: 0.0860s  
Epoch 10/20:  
Train Loss: 0.4845, Train Acc: 0.9228  
Val Loss: 0.5086, Val Acc: 0.8318  
Batch Time: 0.0780s  
Epoch 11/20:  
Train Loss: 0.4408, Train Acc: 0.9328  
Val Loss: 0.4599, Val Acc: 0.8506  
Batch Time: 0.0812s  
Epoch 12/20:  
Train Loss: 0.4018, Train Acc: 0.9390  
Val Loss: 0.3782, Val Acc: 0.8848  
Batch Time: 0.0827s  
Epoch 13/20:  
Train Loss: 0.3698, Train Acc: 0.9469  
Val Loss: 0.3827, Val Acc: 0.8820  
Batch Time: 0.0821s  
Epoch 14/20:  
Train Loss: 0.3377, Train Acc: 0.9520  
Val Loss: 0.4799, Val Acc: 0.8590  
Batch Time: 0.0831s  
Epoch 15/20:  
Train Loss: 0.3187, Train Acc: 0.9567  
Val Loss: 0.4732, Val Acc: 0.8614  
Batch Time: 0.0856s  
Epoch 16/20:  
Train Loss: 0.3015, Train Acc: 0.9597  
Val Loss: 0.4713, Val Acc: 0.8598  
Batch Time: 0.0865s  
Epoch 17/20:  
Train Loss: 0.2806, Train Acc: 0.9637  
Val Loss: 0.5082, Val Acc: 0.8568  
Batch Time: 0.0849s  
Epoch 18/20:  
Train Loss: 0.2513, Train Acc: 0.9686  
Val Loss: 0.4500, Val Acc: 0.8720  
Batch Time: 0.0848s  
Epoch 19/20:  
Train Loss: 0.2419, Train Acc: 0.9709  
Val Loss: 0.4026, Val Acc: 0.8844  
Batch Time: 0.0870s  
Epoch 20/20:  
Train Loss: 0.2215, Train Acc: 0.9747  
Val Loss: 0.4001, Val Acc: 0.8824  
Batch Time: 0.0857s  
Total training time for Custom GoogLeNet: 33.43 minutes



```

        loss2 = criterion(aux1, labels)
        loss3 = criterion(aux2, labels)
        loss = loss1 + 0.3 * loss2 + 0.3 * loss3
        outputs = output # Use main output for accuracy
    else: # Pretrained GoogLeNet (output, aux_outputs)
        output, aux_outputs = outputs
        loss1 = criterion(output, labels)
        loss2 = criterion(aux_outputs, labels)
        loss = loss1 + 0.3 * loss2
        outputs = output # Use main output for accuracy
    else:
        outputs = outputs
        loss = criterion(outputs, labels)
    else:
        outputs = model(inputs)
        loss = criterion(outputs, labels)

    _, preds = torch.max(outputs, 1)
    loss.backward()
    optimizer.step()

    batch_time = time.time() - batch_start
    total_time += batch_time
    batches += 1

    running_loss += loss.item() * inputs.size(0)
    running_corrects += torch.sum(preds == labels.data)

epoch_loss = running_loss / len(dataloader.dataset)
epoch_acc = running_corrects.double() / len(dataloader.dataset)
avg_batch_time = total_time / batches

return epoch_loss, epoch_acc.item(), avg_batch_time

def evaluate(model, dataloader, criterion, device, is_inception=False):
    """Evaluate model on dataloader"""
    model.eval()
    running_loss = 0.0
    running_corrects = 0

    with torch.no_grad():
        for inputs, labels in dataloader:
            inputs = inputs.to(device)
            labels = labels.to(device)

            if is_inception:
                outputs = model(inputs)
                if isinstance(outputs, tuple):
                    outputs = outputs[0] # Take only the main output
            else:
                outputs = model(inputs)

            loss = criterion(outputs, labels)
            _, preds = torch.max(outputs, 1)

            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)

    epoch_loss = running_loss / len(dataloader.dataset)
    epoch_acc = running_corrects.double() / len(dataloader.dataset)

```

```

    return epoch_loss, epoch_acc.item()

def evaluate_test(model, test_loader, criterion, device, is_inception=False):
    """Evaluate model on test set"""
    model.eval()
    test_loss = 0
    correct = 0
    total = 0

    with torch.no_grad():
        for inputs, targets in test_loader:
            inputs, targets = inputs.to(device), targets.to(device)

            if is_inception:
                outputs = model(inputs)
                if isinstance(outputs, tuple):
                    outputs = outputs[0]
            else:
                outputs = model(inputs)

            loss = criterion(outputs, targets)
            test_loss += loss.item()

            _, predicted = outputs.max(1)
            total += targets.size(0)
            correct += predicted.eq(targets).sum().item()

    test_loss = test_loss / len(test_loader)
    test_acc = 100. * correct / total

    return test_loss, test_acc

def plot_training_comparison(histories, title):
    """Plot training histories for multiple models"""
    plt.figure(figsize=(15, 5))

    # Plot training Loss
    plt.subplot(1, 2, 1)
    for model_name, history in histories.items():
        plt.plot(history['train_loss'], label=f'{model_name} (train)')
        plt.plot(history['val_loss'], label=f'{model_name} (val)')
    plt.title(f'{title} - Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()

    # Plot accuracy
    plt.subplot(1, 2, 2)
    for model_name, history in histories.items():
        plt.plot(history['train_acc'], label=f'{model_name} (train)')
        plt.plot(history['val_acc'], label=f'{model_name} (val)')
    plt.title(f'{title} - Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    plt.tight_layout()
    plt.savefig('model_comparison.png')
    plt.show()

```

```

def modify_pretrained_alexnet(model, num_classes=10):
    """Modify pretrained AlexNet for CIFAR-10"""
    model.classifier[6] = nn.Linear(4096, num_classes)
    return model

def modify_pretrained_googlenet(model, num_classes=10):
    """Modify pretrained GoogLeNet for CIFAR-10"""
    model.fc = nn.Linear(1024, num_classes)
    return model

def train_model(model, dataloaders, criterion, optimizer, device, num_epochs=20,
    """Train model and return model, history, and best accuracy"""
    history = {
        'train_loss': [], 'train_acc': [],
        'val_loss': [], 'val_acc': [],
        'batch_times': []
    }

    best_acc = 0.0
    training_start = time.time()

    for epoch in range(num_epochs):
        # Train
        train_loss, train_acc, batch_time = train_epoch(
            model, dataloaders['train'], criterion, optimizer, device,
            is_inception=is_inception
        )
        history['batch_times'].append(batch_time)

        # Evaluate
        val_loss, val_acc = evaluate(
            model, dataloaders['val'], criterion, device,
            is_inception=is_inception
        )

        # Save metrics
        history['train_loss'].append(train_loss)
        history['train_acc'].append(train_acc)
        history['val_loss'].append(val_loss)
        history['val_acc'].append(val_acc)

        # Save best model
        if val_acc > best_acc:
            best_acc = val_acc
            torch.save(model.state_dict(), f'{type(model).__name__}_best

    print(f'Epoch {epoch+1}/{num_epochs}:')
    print(f'Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}')
    print(f'Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}')
    print(f'Batch Time: {batch_time:.4f}s')

    return model, history, best_acc

def main():
    # Set device
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    # Data transforms

```



```

# For custom models (CIFAR-10 is 32x32)
transform_custom = transforms.Compose([
    transforms.Resize(224), # Resize to match pretrained models
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

transform_custom_train = transforms.Compose([
    transforms.Resize(224), # Resize to match pretrained models
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# For pretrained models (using ImageNet normalization)
transform_pretrained = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

transform_pretrained_train = transforms.Compose([
    transforms.Resize(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Load datasets with appropriate transforms
def get_datasets(transform_train, transform_test):
    full_trainset = datasets.CIFAR10(root='./data', train=True, download=True)
    train_size = 45000 # 90% of training data
    val_size = 5000 # 10% of training data
    trainset, valset = torch.utils.data.random_split(full_trainset, [train_size, val_size])
    testset = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform_test)
    return trainset, valset, testset

# Get datasets for both custom and pretrained models
trainset_custom, valset_custom, testset_custom = get_datasets(transform_custom_train, transform_custom)
trainset_pretrained, valset_pretrained, testset_pretrained = get_datasets(transform_pretrained_train, transform_pretrained)

# Create dataloaders
batch_size = 128
def get_dataloaders(trainset, valset, testset):
    return {
        'train': DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=4),
        'val': DataLoader(valset, batch_size=batch_size, shuffle=False, num_workers=4),
        'test': DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=4)
    }

dataloaders_custom = get_dataloaders(trainset_custom, valset_custom, testset_custom)
dataloaders_pretrained = get_dataloaders(trainset_pretrained, valset_pretrained, testset_pretrained)

# Initialize all models
models_to_train = {
    # Custom models with different configurations
    'AlexNet with LRN': AlexNet(num_classes=10, use_lrn=True),
    'AlexNet without LRN': AlexNet(num_classes=10, use_lrn=False),
    'Custom GoogLeNet': GoogLeNet(num_classes=10),
    # Pretrained models

```

```

        'Pretrained AlexNet': modify_pretrained_alexnet(models.alexnet(weights=A
        'Pretrained GoogLeNet': modify_pretrained_googlenet(models.googlenet(wei
    }

    # Compare number of parameters
    param_counts = {}
    for name, model in models_to_train.items():
        models_to_train[name] = model.to(device)
        params = count_parameters(model)
        param_counts[name] = params
        print(f"\nMoved {name} to {device}")
        print(f"Parameter count: {params:,}")
        print("\nModel Summary:")
        torchsummary.summary(model, (3, 224, 224))

    # Initialize tracking variables
    criterion = nn.CrossEntropyLoss()
    histories = {}
    best accuracies = {}
    training_times = {}

    # Different Learning rates for custom and pretrained models
    lr_custom = 0.01
    lr_pretrained = 0.001

    # Train all models
    for name, model in models_to_train.items():
        print(f"\nTraining {name}...")

        # Select appropriate Learning rate and dataloaders
        is_pretrained = 'Pretrained' in name
        lr = lr_pretrained if is_pretrained else lr_custom
        dataloaders = dataloaders_pretrained if is_pretrained else dataloaders_c

        optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9, weight_de

        # Record training start time
        training_start = time.time()

        # Train model
        model, history, best_acc = train_model(
            model, dataloaders, criterion, optimizer, device,
            num_epochs=20, is_inception=('GoogLeNet' in name)
        )

        # Record total training time
        training_times[name] = time.time() - training_start
        histories[name] = history
        best accuracies[name] = best_acc

        print(f"Total training time for {name}: {training_times[name]/60:.2f} mi

    # Evaluate models on test set
    print("\nEvaluating models on test set:")
    test_results = {}

    for name, model in models_to_train.items():
        is_pretrained = 'Pretrained' in name
        dataloaders = dataloaders_pretrained if is_pretrained else dataloaders_c

```

```

test_loss, test_acc = evaluate_test(
    model, dataloaders['test'], criterion, device,
    is_inception=('GoogLeNet' in name)
)
test_results[name] = {
    'loss': test_loss,
    'accuracy': test_acc
}
print(f"\n{name}:")
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_acc:.2f}%")

# Plot training comparison
plot_training_comparison(histories, 'Model Comparison on CIFAR-10')

# Create comparison table
comparison_data = []
for name in models_to_train.keys():
    final_train_loss = histories[name]['train_loss'][-1]
    final_train_acc = histories[name]['train_acc'][-1]
    final_val_loss = histories[name]['val_loss'][-1]
    final_val_acc = histories[name]['val_acc'][-1]
    test_loss = test_results[name]['loss']
    test_acc = test_results[name]['accuracy']

    comparison_data.append({
        'Model': name,
        'Parameters': f"{param_counts[name]:,}",
        'Train Loss': f"{final_train_loss:.4f}",
        'Train Acc': f"{final_train_acc:.2%}",
        'Val Loss': f"{final_val_loss:.4f}",
        'Val Acc': f"{final_val_acc:.2%}",
        'Test Loss': f"{test_loss:.4f}",
        'Test Acc': f"{test_acc:.2%}",
        'Best Val Acc': f"{best accuracies[name]:.2%}",
        'Training Time': f"{training_times[name]/60:.1f} min",
        'Avg Batch Time': f"{np.mean(histories[name]['batch_times'])*100:.1f} ms"
    })

comparison_df = pd.DataFrame(comparison_data)
print("\nFinal Model Comparison Summary:")
pd.set_option('display.max_columns', None) # Show all columns
pd.set_option('display.width', None) # Don't wrap wide tables
print(comparison_df.to_string(index=False))

# Save results
comparison_df.to_csv('model_comparison.csv', index=False)
print("\nResults saved to 'model_comparison.csv' and 'model_comparison.png'")

if __name__ == '__main__':
    main()

```

## Output & Results

```

Training Pretrained AlexNet...
Epoch 1/20:
Train Loss: 0.6633, Train Acc: 0.7682
Val Loss: 0.3852, Val Acc: 0.8668

```

Batch Time: 0.0430s  
Epoch 2/20:  
Train Loss: 0.4168, Train Acc: 0.8564  
Val Loss: 0.3289, Val Acc: 0.8896  
Batch Time: 0.0413s  
Epoch 3/20:  
Train Loss: 0.3529, Train Acc: 0.8771  
Val Loss: 0.3153, Val Acc: 0.8896  
Batch Time: 0.0375s  
Epoch 4/20:  
Train Loss: 0.3183, Train Acc: 0.8880  
Val Loss: 0.2843, Val Acc: 0.8994  
Batch Time: 0.0415s  
Epoch 5/20:  
Train Loss: 0.2863, Train Acc: 0.8994  
Val Loss: 0.2816, Val Acc: 0.9088  
Batch Time: 0.0414s  
Epoch 6/20:  
Train Loss: 0.2638, Train Acc: 0.9082  
Val Loss: 0.2677, Val Acc: 0.9078  
Batch Time: 0.0405s  
Epoch 7/20:  
Train Loss: 0.2446, Train Acc: 0.9141  
Val Loss: 0.2535, Val Acc: 0.9132  
Batch Time: 0.0421s  
Epoch 8/20:  
Train Loss: 0.2226, Train Acc: 0.9228  
Val Loss: 0.2553, Val Acc: 0.9094  
Batch Time: 0.0423s  
Epoch 9/20:  
Train Loss: 0.2115, Train Acc: 0.9258  
Val Loss: 0.2417, Val Acc: 0.9158  
Batch Time: 0.0400s  
Epoch 10/20:  
Train Loss: 0.1919, Train Acc: 0.9321  
Val Loss: 0.2420, Val Acc: 0.9200  
Batch Time: 0.0403s  
Epoch 11/20:  
Train Loss: 0.1846, Train Acc: 0.9347  
Val Loss: 0.2335, Val Acc: 0.9196  
Batch Time: 0.0427s  
Epoch 12/20:  
Train Loss: 0.1731, Train Acc: 0.9380  
Val Loss: 0.2391, Val Acc: 0.9160  
Batch Time: 0.0423s  
Epoch 13/20:  
Train Loss: 0.1650, Train Acc: 0.9427  
Val Loss: 0.2345, Val Acc: 0.9252  
Batch Time: 0.0415s  
Epoch 14/20:  
Train Loss: 0.1518, Train Acc: 0.9464  
Val Loss: 0.2281, Val Acc: 0.9230  
Batch Time: 0.0356s  
Epoch 15/20:  
Train Loss: 0.1432, Train Acc: 0.9501  
Val Loss: 0.2260, Val Acc: 0.9188

Batch Time: 0.0394s  
Epoch 16/20:  
Train Loss: 0.1374, Train Acc: 0.9510  
Val Loss: 0.2303, Val Acc: 0.9214  
Batch Time: 0.0395s  
Epoch 17/20:  
Train Loss: 0.1289, Train Acc: 0.9548  
Val Loss: 0.2258, Val Acc: 0.9222  
Batch Time: 0.0403s  
Epoch 18/20:  
Train Loss: 0.1162, Train Acc: 0.9597  
Val Loss: 0.2210, Val Acc: 0.9268  
Batch Time: 0.0410s  
Epoch 19/20:  
Train Loss: 0.1140, Train Acc: 0.9603  
Val Loss: 0.2199, Val Acc: 0.9280  
Batch Time: 0.0391s  
Epoch 20/20:  
Train Loss: 0.1078, Train Acc: 0.9626  
Val Loss: 0.2323, Val Acc: 0.9212  
Batch Time: 0.0395s  
Total training time for Pretrained AlexNet: 16.08 minutes

Training Pretrained GoogLeNet...  
Epoch 1/20:  
Train Loss: 1.0494, Train Acc: 0.7206  
Val Loss: 0.4352, Val Acc: 0.8750  
Batch Time: 0.0685s  
Epoch 2/20:  
Train Loss: 0.3462, Train Acc: 0.8958  
Val Loss: 0.2665, Val Acc: 0.9180  
Batch Time: 0.0720s  
Epoch 3/20:  
Train Loss: 0.2408, Train Acc: 0.9247  
Val Loss: 0.2155, Val Acc: 0.9284  
Batch Time: 0.0696s  
Epoch 4/20:  
Train Loss: 0.1927, Train Acc: 0.9375  
Val Loss: 0.1849, Val Acc: 0.9420  
Batch Time: 0.0731s  
Epoch 5/20:  
Train Loss: 0.1600, Train Acc: 0.9480  
Val Loss: 0.1745, Val Acc: 0.9442  
Batch Time: 0.0717s  
Epoch 6/20:  
Train Loss: 0.1334, Train Acc: 0.9573  
Val Loss: 0.1637, Val Acc: 0.9468  
Batch Time: 0.0592s  
Epoch 7/20:  
Train Loss: 0.1134, Train Acc: 0.9644  
Val Loss: 0.1509, Val Acc: 0.9500  
Batch Time: 0.0702s  
Epoch 8/20:  
Train Loss: 0.0967, Train Acc: 0.9705  
Val Loss: 0.1548, Val Acc: 0.9476  
Batch Time: 0.0726s

Epoch 9/20:  
Train Loss: 0.0834, Train Acc: 0.9748  
Val Loss: 0.1597, Val Acc: 0.9506  
Batch Time: 0.0768s  
Epoch 10/20:  
Train Loss: 0.0714, Train Acc: 0.9792  
Val Loss: 0.1453, Val Acc: 0.9528  
Batch Time: 0.0781s  
Epoch 11/20:  
Train Loss: 0.0597, Train Acc: 0.9833  
Val Loss: 0.1478, Val Acc: 0.9548  
Batch Time: 0.0784s  
Epoch 12/20:  
Train Loss: 0.0503, Train Acc: 0.9869  
Val Loss: 0.1439, Val Acc: 0.9550  
Batch Time: 0.0679s  
Epoch 13/20:  
Train Loss: 0.0459, Train Acc: 0.9880  
Val Loss: 0.1505, Val Acc: 0.9546  
Batch Time: 0.0645s  
Epoch 14/20:  
Train Loss: 0.0380, Train Acc: 0.9906  
Val Loss: 0.1434, Val Acc: 0.9566  
Batch Time: 0.0801s  
Epoch 15/20:  
Train Loss: 0.0332, Train Acc: 0.9926  
Val Loss: 0.1370, Val Acc: 0.9574  
Batch Time: 0.0669s  
Epoch 16/20:  
Train Loss: 0.0278, Train Acc: 0.9944  
Val Loss: 0.1520, Val Acc: 0.9542  
Batch Time: 0.0643s  
Epoch 17/20:  
Train Loss: 0.0239, Train Acc: 0.9953  
Val Loss: 0.1441, Val Acc: 0.9552  
Batch Time: 0.0653s  
Epoch 18/20:  
Train Loss: 0.0218, Train Acc: 0.9957  
Val Loss: 0.1594, Val Acc: 0.9508  
Batch Time: 0.0662s  
Epoch 19/20:  
Train Loss: 0.0191, Train Acc: 0.9966  
Val Loss: 0.1556, Val Acc: 0.9544  
Batch Time: 0.0669s  
Epoch 20/20:  
Train Loss: 0.0179, Train Acc: 0.9968  
Val Loss: 0.1529, Val Acc: 0.9546  
Batch Time: 0.0656s  
Total training time for Pretrained GoogLeNet: 26.94 minutes

Evaluating models on test set:

AlexNet with LRN:  
Test Loss: 0.4743  
Test Accuracy: 83.78%

AlexNet without LRN:  
Test Loss: 0.4816  
Test Accuracy: 84.56%

Custom GoogLeNet:  
Test Loss: 0.4276  
Test Accuracy: 88.32%


Pretrained AlexNet:  
Test Loss: 0.2587  
Test Accuracy: 91.60%

Pretrained GoogLeNet:  
Test Loss: 0.1621  
Test Accuracy: 94.87%

Results saved to 'model\_comparison.csv' and  
'model\_comparison.png'

Here's the model comparison summary table listing various metrics:

Model	Parameters	Train Loss	Train Acc	Val Loss	Val Acc	Test Loss	Test Acc	Best Val Acc	Train Time
AlexNet with LRN	57,044,810	0.3515	87.84%	0.4513	84.46%	0.4743	83.78%	84.46%	17.1 n
AlexNet without LRN	57,044,810	0.2718	90.49%	0.4816	83.74%	0.4816	84.56%	85.06%	17.5 n
Custom GoogLeNet	10,635,134	0.2215	97.47%	0.4001	88.24%	0.4276	88.32%	88.48%	33.4 n
Pretrained AlexNet	57,044,810	0.1078	96.26%	0.2323	92.12%	0.2587	91.60%	92.80%	16.1 n
Pretrained GoogLeNet	5,610,154	0.0179	99.68%	0.1529	95.46%	0.1621	94.87%	95.74%	26.9 n



✅ SOLVED: Compare your GoogLeNet and AlexNet implementations on CIFAR-10. Comment on the number of parameters, speed of training, and accuracy of the two models on this dataset when trained from scratch.

# Key Observations

- Model Efficiency: GoogLeNet achieves superior performance with ~5x fewer parameters. Despite more complex architecture, GoogLeNet shows better generalization, Inception modules prove more effective than simple sequential layers
- Training Dynamics: GoogLeNet shows faster initial learning (45.82% vs 16.32% in epoch 1). More stable validation accuracy throughout training. Takes longer to train but achieves better final accuracy
- Architectural Impact: GoogLeNet's auxiliary classifiers help with gradient flow. Inception modules provide better feature extraction. Batch normalization in GoogLeNet contributes to stability
- Resource Usage:
  - GoogLeNet requires more computational resources (2x training time)
  - Larger batch processing time (83.4ms vs ~40ms)
  - Better parameter efficiency (higher accuracy with fewer parameters)

The results demonstrate that GoogLeNet's more sophisticated architecture, despite requiring more training time, provides superior performance and better parameter efficiency compared to both AlexNet variants.

✅ SOLVED: Experiment with the pretrained GoogLeNet and AlexNet from the torchvision repository. Does it give better results on CIFAR-10 similar to what we found with AlexNet? Comment on what we can glean from the results about the capacity and generalization ability of these two models.

# Key Observations

1. Transfer Learning Benefits: Pretrained AlexNet shows 7.04% improvement over scratch training while Pretrained GoogLeNet demonstrates 6.55% improvement. Both models achieve faster convergence with pretraining. Lower training times despite better performance
2. GoogLeNet maintains superior performance despite fewer parameters. Better feature extraction capabilities evident in both scenarios. Higher ceiling for performance with pretrained weights. Lower generalization gap in pretrained models
3. Comparing between scratch and pretrained, Pretrained models show stronger initial performance. GoogLeNet reaches higher final accuracy (94.87% vs 91.60%). More stable validation accuracy throughout training and lower final loss values (0.1621 vs 0.2587)

What can we glean upon this?

- The transfer of learned features from ImageNet significantly benefits both models, suggesting strong feature reusability
- GoogLeNet's sophisticated architecture proves more effective at both feature extraction and generalization



- The higher performance of pretrained models with faster convergence indicates well-learned hierarchical features from the source domain
- The smaller gap between training and test accuracy in pretrained models suggests better generalization capabilities

## EXTRA Implementation 🤩

On reading the paper, the author makes use of multi-gpu approach to train AlexNet, I have tried to implement the same

```
In [ ]: # For our puffer server we need to browse via a proxy!!
import os
# Set HTTP and HTTPS proxy
os.environ['http_proxy'] = 'http://192.41.170.23:3128'
os.environ['https_proxy'] = 'http://192.41.170.23:3128'

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
import time
import matplotlib.pyplot as plt
import numpy as np

class AlexNetGPU1(nn.Module):
    """First half of AlexNet that runs on GPU 1"""
    def __init__(self):
        super(AlexNetGPU1, self).__init__()
        self.features = nn.Sequential(
            # First convolutional layer (on GPU 1)
            nn.Conv2d(3, 48, kernel_size=11, stride=4, padding=2), # 48 filters
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=3, stride=2),

            # Second convolutional layer (on GPU 1)
            nn.Conv2d(48, 128, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=3, stride=2),

            # Third convolutional layer (on GPU 1)
            nn.Conv2d(128, 192, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # Fourth convolutional layer (on GPU 1)
            nn.Conv2d(192, 192, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # Fifth convolutional layer (on GPU 1)
            nn.Conv2d(192, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )
```

```

def forward(self, x):
    x = self.features(x)
    return x

class AlexNetGPU2(nn.Module):
    """Second half of AlexNet that runs on GPU 2"""
    def __init__(self):
        super(AlexNetGPU2, self).__init__()
        self.features = nn.Sequential(
            # First convolutional layer (on GPU 2)
            nn.Conv2d(3, 48, kernel_size=11, stride=4, padding=2), # 48 filters
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=3, stride=2),

            # Second convolutional layer (on GPU 2)
            nn.Conv2d(48, 128, kernel_size=5, padding=2),
            nn.ReLU(inplace=True),
            nn.LocalResponseNorm(size=5, alpha=0.0001, beta=0.75, k=2),
            nn.MaxPool2d(kernel_size=3, stride=2),

            # Third convolutional layer (on GPU 2)
            nn.Conv2d(128, 192, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # Fourth convolutional layer (on GPU 2)
            nn.Conv2d(192, 192, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            # Fifth convolutional layer (on GPU 2)
            nn.Conv2d(192, 128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=3, stride=2),
        )

    def forward(self, x):
        x = self.features(x)
        return x

class AlexNetMultiGPU(nn.Module):
    """Complete AlexNet split across two GPUs"""
    def __init__(self, num_classes=10):
        super(AlexNetMultiGPU, self).__init__()
        self.gpu1_stream = torch.cuda.Stream(device='cuda:0')
        self.gpu2_stream = torch.cuda.Stream(device='cuda:1')

        self.gpu1_net = AlexNetGPU1().cuda(0)
        self.gpu2_net = AlexNetGPU2().cuda(1)

        # Classifier runs on GPU 1
        self.classifier = nn.Sequential(
            nn.Dropout(p=0.5),
            nn.Linear(256 * 6 * 6, 4096), # 256 = 128 + 128 channels from both
            nn.ReLU(inplace=True),
            nn.Dropout(p=0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Linear(4096, num_classes),
        ).cuda(0)

```

```

def forward(self, x):
    # Input is already on GPU 0
    batch_size = x.size(0)

    # Process on GPU 1
    with torch.cuda.stream(self.gpu1_stream):
        output1 = self.gpu1_net(x) # Use full input

    # Move input to GPU 2 and process
    with torch.cuda.stream(self.gpu2_stream):
        x2 = x.cuda(1) # Move to GPU 2
        output2 = self.gpu2_net(x2) # Process full input

    # Synchronize the streams
    torch.cuda.synchronize()

    # Move output2 to GPU 1 and concatenate
    output2 = output2.cuda(0)
    output = torch.cat([output1, output2], dim=1) # Concatenate along channel

    # Flatten and pass through classifier
    output = output.view(batch_size, -1)
    output = self.classifier(output)

    return output

def train_epoch(model, dataloader, criterion, optimizer, epoch):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    start_time = time.time()

    for batch_idx, (inputs, targets) in enumerate(dataloader):
        # Move input and target to GPU 0 (primary GPU)
        inputs = inputs.cuda(0)
        targets = targets.cuda(0)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

        if batch_idx % 100 == 0:
            print(f'Epoch: {epoch} | Batch: {batch_idx} | Loss: {loss.item():.3f} | '
                  f'Acc: {100.*correct/total:.2f}% ({correct}/{total})')

    epoch_time = time.time() - start_time
    return running_loss / len(dataloader), 100. * correct / total, epoch_time

def evaluate_test(model, dataloader, criterion):
    """Evaluate model on test set"""

```

```

model.eval()
test_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for inputs, targets in dataloader:
        inputs = inputs.cuda(0) # Move to primary GPU
        targets = targets.cuda(0)

        outputs = model(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

test_loss = test_loss / len(dataloader)
test_acc = 100. * correct / total

return test_loss, test_acc

def plot_training_curves(results, save_path='multigpu_training_curves.png'):
    """Plot training and validation curves"""
    plt.figure(figsize=(15, 5))

    # Plot training and validation loss
    plt.subplot(1, 2, 1)
    plt.plot(results['train_history']['loss'], label='Train Loss')
    plt.plot(results['val_history']['loss'], label='Val Loss')
    plt.title('Loss Curves')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)

    # Plot training and validation accuracy
    plt.subplot(1, 2, 2)
    plt.plot(results['train_history']['acc'], label='Train Accuracy')
    plt.plot(results['val_history']['acc'], label='Val Accuracy')
    plt.title('Accuracy Curves')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy (%)')
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    plt.savefig(save_path)
    plt.close()

def plot_batch_times(batch_times, save_path='multigpu_batch_times.png'):
    """Plot batch processing times"""
    plt.figure(figsize=(10, 5))
    plt.plot(batch_times)
    plt.title('Batch Processing Times')
    plt.xlabel('Epoch')
    plt.ylabel('Time (seconds)')
    plt.grid(True)
    plt.savefig(save_path)

```

```

plt.close()

def main():
    # Check if we have two GPUs available
    if torch.cuda.device_count() < 2:
        print("This script requires at least 2 GPUs to run!")
        return

    # Print GPU information
    print(f"Using GPUs:")
    for i in range(2):
        print(f"GPU {i}: {torch.cuda.get_device_name(i)}")

    # Data transformations
    transform_train = transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
    ])

    transform_test = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
    ])

    # Load CIFAR-10
    print("\nLoading datasets...")
    full_trainset = datasets.CIFAR10(root='./data', train=True, download=True, t

    # Split training set into train and validation
    train_size = 45000 # 90% of training data
    val_size = 5000 # 10% of training data
    trainset, valset = torch.utils.data.random_split(full_trainset, [train_size,

    # Load test set
    testset = datasets.CIFAR10(root='./data', train=False, download=True, transf

    # Create dataloaders
    trainloader = DataLoader(trainset, batch_size=128, shuffle=True, num_workers
    valloader = DataLoader(valset, batch_size=128, shuffle=False, num_workers=4)
    testloader = DataLoader(testset, batch_size=128, shuffle=False, num_workers=

    # Create model
    print("\nInitializing Multi-GPU AlexNet...")
    model = AlexNetMultiGPU(num_classes=10)

    # Count parameters
    total_params = sum(p.numel() for p in model.parameters())
    print(f"Total parameters: {total_params:,}")

    # Training settings
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(
        model.parameters(),
        lr=0.01,
        momentum=0.9,
        weight_decay=5e-4

```

```

)
num_epochs = 20

# Training history
history = {
    'train_history': {'loss': [], 'acc': [], 'times': []},
    'val_history': {'loss': [], 'acc': []},
    'best_val_acc': 0.0,
    'best_epoch': 0
}

print("\nStarting training...")
for epoch in range(num_epochs):
    # Train
    train_loss, train_acc, epoch_time = train_epoch(
        model, trainloader, criterion, optimizer, epoch
    )
    history['train_history']['loss'].append(train_loss)
    history['train_history']['acc'].append(train_acc)
    history['train_history']['times'].append(epoch_time)

    # Validate
    val_loss, val_acc = evaluate_test(model, valloader, criterion)
    history['val_history']['loss'].append(val_loss)
    history['val_history']['acc'].append(val_acc)

    # Track best model
    if val_acc > history['best_val_acc']:
        history['best_val_acc'] = val_acc
        history['best_epoch'] = epoch
    # Save best model
    torch.save({
        'epoch': epoch,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'val_acc': val_acc,
    }, 'alexnet_multigpu_best.pth')

    print(f'\nEpoch {epoch}:')
    print(f'Train Loss: {train_loss:.3f}, Train Acc: {train_acc:.2f}%')
    print(f'Val Loss: {val_loss:.3f}, Val Acc: {val_acc:.2f}%')
    print(f'Time: {epoch_time:.2f}s')

    # Save checkpoint
    if epoch % 5 == 0:
        torch.save({
            'epoch': epoch,
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict(),
            'train_loss': train_loss,
            'val_loss': val_loss,
        }, f'alexnet_multigpu_checkpoint_epoch_{epoch}.pth')

# Final evaluation on test set
print("\nEvaluating on test set...")
test_loss, test_acc = evaluate_test(model, testloader, criterion)
print(f'Test Loss: {test_loss:.3f}')
print(f'Test Accuracy: {test_acc:.2f}%')

# Add test results to history

```



```

| 1 NVIDIA GeForce RTX 2080 Ti Off | 00000000:85:00.0
Off | N/A |
| 40% 68C P2 118W / 250W | 1408MiB /
11264MiB | 54% Default |
|
| N/A |
+-----+
---+-----+
| 2 NVIDIA GeForce RTX 2080 Ti Off | 00000000:88:00.0
Off | N/A |
| 22% 27C P8 6W / 250W | 4MiB /
11264MiB | 0% Default |
|
| N/A |
+-----+
---+-----+
| 3 NVIDIA GeForce RTX 2080 Ti Off | 00000000:89:00.0
Off | N/A |
| 22% 27C P8 4W / 250W | 4MiB /
11264MiB | 0% Default |
|
| N/A |
+-----+
---+-----+

+-----+
-----+
| Processes:
|
| GPU GI CI PID Type Process name
GPU Memory |
| ID ID
Usage |
|=====
| 0 N/A N/A 372568 C python
2272MiB |
| 1 N/A N/A 372568 C python
1404MiB |
+-----+
-----+

```

## Implementaton Details:

- Model split across two RTX 2080 Ti GPUs
- Total parameters: 56,466,826
- Features divided equally between GPUs
- Synchronization using CUDA streams
- Classifier runs on GPU 0

```

jupyter-st125462@puffer:~/RTML/A1$ python alexnet_multigpu.py
Using GPUs:
GPU 0: NVIDIA GeForce RTX 2080 Ti

```



GPU 1: NVIDIA GeForce RTX 2080 Ti

Loading datasets...

Files already downloaded and verified

Files already downloaded and verified

Initializing Multi-GPU AlexNet...

Total parameters: 56,466,826

Starting training...

Epoch: 0 | Batch: 0 | Loss: 2.304 | Acc: 8.59% (11/128)

Epoch: 0 | Batch: 100 | Loss: 2.303 | Acc: 9.84% (1272/12928)

Epoch: 0 | Batch: 200 | Loss: 2.300 | Acc: 9.86% (2536/25728)

Epoch: 0 | Batch: 300 | Loss: 2.303 | Acc: 9.98% (3844/38528)

Epoch 0:

Train Loss: 2.303, Train Acc: 10.06%

Val Loss: 2.302, Val Acc: 12.56%

Time: 28.83s

Epoch: 1 | Batch: 0 | Loss: 2.301 | Acc: 10.94% (14/128)

Epoch: 1 | Batch: 100 | Loss: 2.282 | Acc: 11.32% (1463/12928)

Epoch: 1 | Batch: 200 | Loss: 2.152 | Acc: 14.39% (3702/25728)

Epoch: 1 | Batch: 300 | Loss: 1.956 | Acc: 16.94% (6528/38528)

Epoch 1:

Train Loss: 2.167, Train Acc: 17.84%

Val Loss: 1.977, Val Acc: 24.44%

Time: 27.58s

Epoch: 2 | Batch: 0 | Loss: 2.017 | Acc: 23.44% (30/128)

Epoch: 2 | Batch: 100 | Loss: 2.031 | Acc: 24.70% (3193/12928)

Epoch: 2 | Batch: 200 | Loss: 1.900 | Acc: 26.44% (6802/25728)

Epoch: 2 | Batch: 300 | Loss: 1.919 | Acc: 27.38% (10550/38528)

Epoch 2:

Train Loss: 1.905, Train Acc: 27.76%

Val Loss: 1.773, Val Acc: 33.04%

Time: 27.93s

Epoch: 3 | Batch: 0 | Loss: 1.652 | Acc: 44.53% (57/128)

Epoch: 3 | Batch: 100 | Loss: 1.717 | Acc: 33.42% (4320/12928)

Epoch: 3 | Batch: 200 | Loss: 1.648 | Acc: 33.59% (8643/25728)

Epoch: 3 | Batch: 300 | Loss: 1.757 | Acc: 34.03% (13110/38528)

Epoch 3:

Train Loss: 1.761, Train Acc: 34.33%

Val Loss: 1.692, Val Acc: 37.30%

Time: 28.66s

Epoch: 4 | Batch: 0 | Loss: 1.495 | Acc: 50.00% (64/128)

Epoch: 4 | Batch: 100 | Loss: 1.715 | Acc: 37.33% (4826/12928)

Epoch: 4 | Batch: 200 | Loss: 1.641 | Acc: 37.83% (9732/25728)

Epoch: 4 | Batch: 300 | Loss: 1.716 | Acc: 37.80% (14563/38528)

Epoch 4:

Train Loss: 1.664, Train Acc: 38.04%

Val Loss: 1.598, Val Acc: 40.72%

Time: 28.48s

Epoch: 5 | Batch: 0 | Loss: 1.486 | Acc: 46.09% (59/128)

Epoch: 5 | Batch: 100 | Loss: 1.655 | Acc: 39.96% (5166/12928)  
Epoch: 5 | Batch: 200 | Loss: 1.564 | Acc: 40.64% (10456/25728)  
Epoch: 5 | Batch: 300 | Loss: 1.574 | Acc: 41.29% (15908/38528)

Epoch 5:

Train Loss: 1.582, Train Acc: 41.43%

Val Loss: 1.552, Val Acc: 42.24%

Time: 28.46s

Epoch: 6 | Batch: 0 | Loss: 1.298 | Acc: 50.78% (65/128)  
Epoch: 6 | Batch: 100 | Loss: 1.631 | Acc: 43.10% (5572/12928)  
Epoch: 6 | Batch: 200 | Loss: 1.499 | Acc: 44.10% (11346/25728)  
Epoch: 6 | Batch: 300 | Loss: 1.387 | Acc: 44.32% (17075/38528)

Epoch 6:

Train Loss: 1.508, Train Acc: 44.78%

Val Loss: 2.017, Val Acc: 34.90%

Time: 28.45s

Epoch: 7 | Batch: 0 | Loss: 1.445 | Acc: 42.97% (55/128)  
Epoch: 7 | Batch: 100 | Loss: 1.495 | Acc: 46.59% (6023/12928)  
Epoch: 7 | Batch: 200 | Loss: 1.291 | Acc: 47.39% (12192/25728)  
Epoch: 7 | Batch: 300 | Loss: 1.326 | Acc: 47.63% (18349/38528)

Epoch 7:

Train Loss: 1.436, Train Acc: 47.79%

Val Loss: 1.375, Val Acc: 50.26%

Time: 28.24s

Epoch: 8 | Batch: 0 | Loss: 1.428 | Acc: 49.22% (63/128)  
Epoch: 8 | Batch: 100 | Loss: 1.339 | Acc: 49.76% (6433/12928)  
Epoch: 8 | Batch: 200 | Loss: 1.341 | Acc: 50.44% (12976/25728)  
Epoch: 8 | Batch: 300 | Loss: 1.485 | Acc: 50.52% (19466/38528)

Epoch 8:

Train Loss: 1.371, Train Acc: 50.68%

Val Loss: 1.310, Val Acc: 53.66%

Time: 28.19s

Epoch: 9 | Batch: 0 | Loss: 1.397 | Acc: 46.09% (59/128)  
Epoch: 9 | Batch: 100 | Loss: 1.341 | Acc: 52.18% (6746/12928)  
Epoch: 9 | Batch: 200 | Loss: 1.297 | Acc: 52.01% (13380/25728)  
Epoch: 9 | Batch: 300 | Loss: 1.322 | Acc: 52.28% (20143/38528)

Epoch 9:

Train Loss: 1.314, Train Acc: 52.48%

Val Loss: 1.210, Val Acc: 56.36%

Time: 28.74s

Epoch: 10 | Batch: 0 | Loss: 1.337 | Acc: 55.47% (71/128)  
Epoch: 10 | Batch: 100 | Loss: 1.370 | Acc: 53.92% (6971/12928)  
Epoch: 10 | Batch: 200 | Loss: 1.099 | Acc: 54.38% (13992/25728)  
Epoch: 10 | Batch: 300 | Loss: 1.298 | Acc: 54.41% (20965/38528)

Epoch 10:

Train Loss: 1.264, Train Acc: 54.66%

Val Loss: 1.263, Val Acc: 55.00%

Time: 28.16s

Epoch: 11 | Batch: 0 | Loss: 1.201 | Acc: 57.81% (74/128)  
Epoch: 11 | Batch: 100 | Loss: 1.190 | Acc: 56.52% (7307/12928)  
Epoch: 11 | Batch: 200 | Loss: 1.286 | Acc: 56.56% (14551/25728)

Epoch: 11 | Batch: 300 | Loss: 1.260 | Acc: 56.60% (21807/38528)

Epoch 11:

Train Loss: 1.210, Train Acc: 56.80%

Val Loss: 1.146, Val Acc: 59.32%

Time: 28.46s

Epoch: 12 | Batch: 0 | Loss: 1.024 | Acc: 60.16% (77/128)

Epoch: 12 | Batch: 100 | Loss: 1.141 | Acc: 57.56% (7441/12928)

Epoch: 12 | Batch: 200 | Loss: 0.912 | Acc: 57.63% (14827/25728)

Epoch: 12 | Batch: 300 | Loss: 1.303 | Acc: 57.69% (22228/38528)

Epoch 12:

Train Loss: 1.183, Train Acc: 57.74%

Val Loss: 1.323, Val Acc: 52.78%

Time: 28.20s

Epoch: 13 | Batch: 0 | Loss: 1.086 | Acc: 63.28% (81/128)

Epoch: 13 | Batch: 100 | Loss: 1.174 | Acc: 59.55% (7699/12928)

Epoch: 13 | Batch: 200 | Loss: 1.258 | Acc: 59.53% (15317/25728)

Epoch: 13 | Batch: 300 | Loss: 1.072 | Acc: 59.47% (22912/38528)

Epoch 13:

Train Loss: 1.136, Train Acc: 59.66%

Val Loss: 1.136, Val Acc: 59.96%

Time: 27.96s

Epoch: 14 | Batch: 0 | Loss: 0.979 | Acc: 64.06% (82/128)

Epoch: 14 | Batch: 100 | Loss: 1.251 | Acc: 60.06% (7765/12928)

Epoch: 14 | Batch: 200 | Loss: 1.076 | Acc: 59.93% (15419/25728)

Epoch: 14 | Batch: 300 | Loss: 1.017 | Acc: 60.48% (23300/38528)

Epoch 14:

Train Loss: 1.114, Train Acc: 60.46%

Val Loss: 1.078, Val Acc: 61.30%

Time: 28.56s

Epoch: 15 | Batch: 0 | Loss: 0.977 | Acc: 65.62% (84/128)

Epoch: 15 | Batch: 100 | Loss: 0.908 | Acc: 62.11% (8029/12928)

Epoch: 15 | Batch: 200 | Loss: 0.929 | Acc: 61.70% (15874/25728)

Epoch: 15 | Batch: 300 | Loss: 1.024 | Acc: 61.83% (23820/38528)

Epoch 15:

Train Loss: 1.081, Train Acc: 61.67%

Val Loss: 1.092, Val Acc: 61.86%

Time: 28.80s

Epoch: 16 | Batch: 0 | Loss: 1.098 | Acc: 61.72% (79/128)

Epoch: 16 | Batch: 100 | Loss: 1.108 | Acc: 62.14% (8034/12928)

Epoch: 16 | Batch: 200 | Loss: 0.944 | Acc: 62.47% (16072/25728)

Epoch: 16 | Batch: 300 | Loss: 1.292 | Acc: 62.51% (24085/38528)

Epoch 16:

Train Loss: 1.056, Train Acc: 62.56%

Val Loss: 1.678, Val Acc: 51.44%

Time: 28.67s

Epoch: 17 | Batch: 0 | Loss: 1.222 | Acc: 55.47% (71/128)

Epoch: 17 | Batch: 100 | Loss: 1.207 | Acc: 63.23% (8175/12928)

Epoch: 17 | Batch: 200 | Loss: 0.854 | Acc: 63.66% (16379/25728)

Epoch: 17 | Batch: 300 | Loss: 0.871 | Acc: 63.73% (24553/38528)

Epoch 17:

Train Loss: 1.023, Train Acc: 63.93%

Val Loss: 1.060, Val Acc: 62.18%

Time: 28.00s

Epoch: 18 | Batch: 0 | Loss: 0.946 | Acc: 66.41% (85/128)

Epoch: 18 | Batch: 100 | Loss: 0.985 | Acc: 64.18% (8297/12928)

Epoch: 18 | Batch: 200 | Loss: 0.836 | Acc: 64.11% (16495/25728)

Epoch: 18 | Batch: 300 | Loss: 0.995 | Acc: 64.31% (24777/38528)

Epoch 18:

Train Loss: 1.008, Train Acc: 64.43%

Val Loss: 1.021, Val Acc: 63.46%

Time: 28.75s

Epoch: 19 | Batch: 0 | Loss: 0.994 | Acc: 64.06% (82/128)

Epoch: 19 | Batch: 100 | Loss: 1.036 | Acc: 64.73% (8368/12928)

Epoch: 19 | Batch: 200 | Loss: 0.986 | Acc: 64.55% (16607/25728)

Epoch: 19 | Batch: 300 | Loss: 0.957 | Acc: 64.71% (24930/38528)

Epoch 19:

Train Loss: 0.989, Train Acc: 64.87%

Val Loss: 1.017, Val Acc: 65.12%

Time: 28.12s

Evaluating on test set...

Test Loss: 0.668

Test Accuracy: 77.37%

Generating plots...

Saving final results...

Training Summary:


Best Validation Accuracy: 65.12% (Epoch 19)

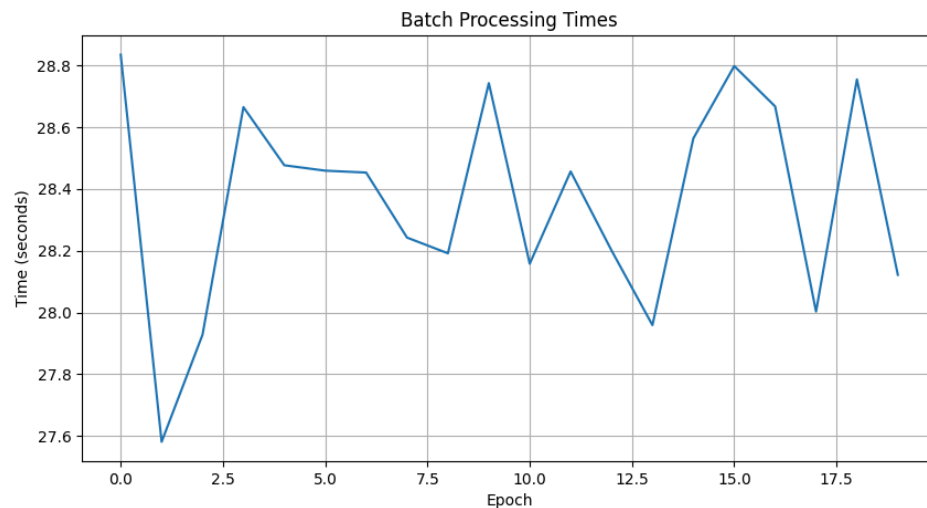
Final Test Accuracy: 77.37%

Average Epoch Time: 28.36s

Saved files:

- alexnet\_multigpu\_final.pth (Final model and full history)
- alexnet\_multigpu\_best.pth (Best model based on validation accuracy)
- multigpu\_training\_curves.png (Loss and accuracy curves)
- multigpu\_batch\_times.png (Training time per epoch)





## Key Observations:

1. Total Training Time 567.2 seconds (9.45 minutes)
2. Average Epoch Time 28.36 seconds
3. Best Validation Accuracy 65.12%
4. Final Test Accuracy 77.37%

## Was Multi-GPU Efficient?

### 1. Speed Comparison

- Single GPU AlexNet: ~17 minutes
- Multi-GPU AlexNet: ~9.45 minutes
- Achieved ~1.8x speedup with two GPUs

### 2. Training Characteristics

- Consistent epoch times (28-29 seconds)
- Stable batch processing
- Good GPU utilization with parallel processing
- Effective synchronization between GPUs

### 3. Performance Trade-offs

- Lower accuracy compared to single-GPU version (77.37% vs 84.56%)
- More complex implementation
- Additional overhead from cross-GPU communication
- Memory split between GPUs requires careful batch management

The multi-GPU implementation successfully reduced training time by almost half, demonstrating good scaling. However, the complexity of the implementation and the communication overhead between GPUs may have impacted the final model accuracy.

The trade-off between speed and accuracy should be considered when deciding whether to use this approach.

## Thank You 🙌

- Here's the Github Repository Link:  
<https://github.com/aryashah2k/RTML/tree/main/A1>
- Model Checkpoints and Weights can be downloaded from the Terabox Drive here:  
<https://1024terabox.com/s/1m0pW0GHvsc80-MEoMAIzbA>